

Process Management

Synchronization Primitives



Lecture 6



<https://archipelago.rocks/app/resend-invite/04498013008>

Overview

■ Race Condition

- ❑ Problems with concurrent execution

■ Critical Section

- ❑ Properties of correct implementation
- ❑ Symptoms of incorrect implementation

■ Implementations of Critical Section


- ❑ High-level programming language solution
- ❑ Low level (hardware) solution
- ❑ High-level OS abstraction

■ Classical synchronization problems (next lecture)

Problems with Concurrent Execution

- Execution of a single sequential process is **deterministic**
 - Repeated execution gives the same result
- When two or more processes:
 - Execute concurrently in interleaving fashion AND
 - Share a modifiable resource
 - Can cause **synchronization problems**
- Execution of concurrent processes may be **non-deterministic**
 - Execution outcome depends on the order in which the shared resource is accessed/modified
 - These problems are known as **race conditions** or **data races**

Race Condition: Illustration



- Process **P1** and **P2** shares a variable **X**
- Rough translation for **X = X + 1000**
 1. **Load** **X** \rightarrow Register1
 2. **Add** 1000 to Register1
 3. **Store** Register1 \rightarrow **X**

Race Condition: Round 1

| Time | Value of X | P1 | P2 |
|------|------------|----|----|
| 1 | 345 | | |

Race Condition: Round 2

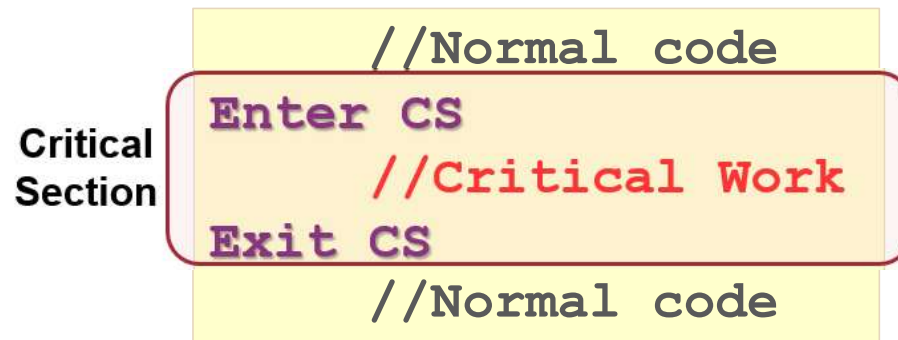
| Time | Value of X | P1 | P2 |
|------|------------|------------------|----|
| 1 | 345 | Load x → Reg1 | |
| 2 | 345 | Add 1000 to Reg1 | |

Race Condition: **Solution**

Incorrect execution is due to the
**unsynchronized access to a
shared modifiable resource**

- Needed: synchronization to control the interleaving of accesses to a shared resource
 - ❑ Allow all (or as many as possible) correct interleaving scenarios
 - ❑ Not allow any incorrect interleaving scenario
- Solution outline:
 - ❑ Designate code segment with race condition as **critical section**
 - ❑ At any point in time, at most **one process** can be in the critical section

Critical Section (CS)



Generic Skeleton of code with Critical Section(s):

Example:

```
Enter CS
    X = X + 1000
Exit CS
```

Process **P1**

```
Enter CS
    X = X + 1000
Exit CS
```

Process **P2**

Race Condition Solution: "Illustration"

Mutual
Exclusion:

Progress:

Bounded Wait:

Independence:



Properties of Correct CS Implementation

Mutual Exclusion:

- If process P_i is executing in critical section, all other processes are prevented from entering the critical section.

Progress:

- If no process is in a critical section, one of the waiting processes should be granted access.

Bounded Wait:

- After process P_i request to enter critical section, there exists an upper bound on the number of times other processes can enter the critical section before P_i .

Independence:

- Process **not** executing in critical section should never block other process.

Symptoms of **Incorrect Synchronization**

■ **Incorrect output/behavior**

- ❑ Usually due to lack of mutual exclusion

■ **Deadlock:**

- ❑ All processes blocked → *no progress*


■ **Livelock:**

- ❑ Typically, processes are not in a blocked state
- ❑ Processes keep changing state to avoid deadlock but make no other progress
- ❑ Usually related to *deadlock avoidance* mechanisms


■ **Starvation:**

- ❑ Some processes are blocked forever

Deadlock



Livelock



CS Implementations Overview

- **High-level programming language implementations:**

- Using only normal programming constructs

- **Assembly-level implementations:**

- Mechanisms provided by the hardware
 - through special instructions

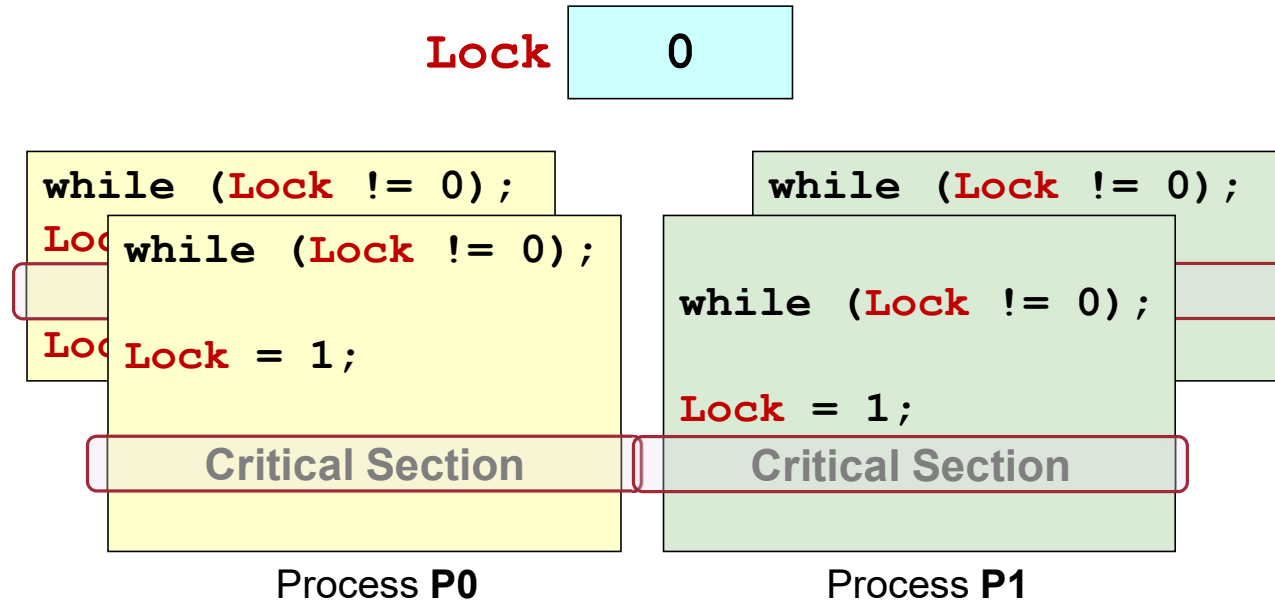
- **High level abstractions:**

- Provide abstracted mechanisms that provide additional useful features
- Commonly implemented by assembly level mechanisms

Using only your brain power.... 😊

HIGH LEVEL LANGUAGE IMPLEMENTATION

Using HLL: **Attempt 1**



Mutual
Exclusion:

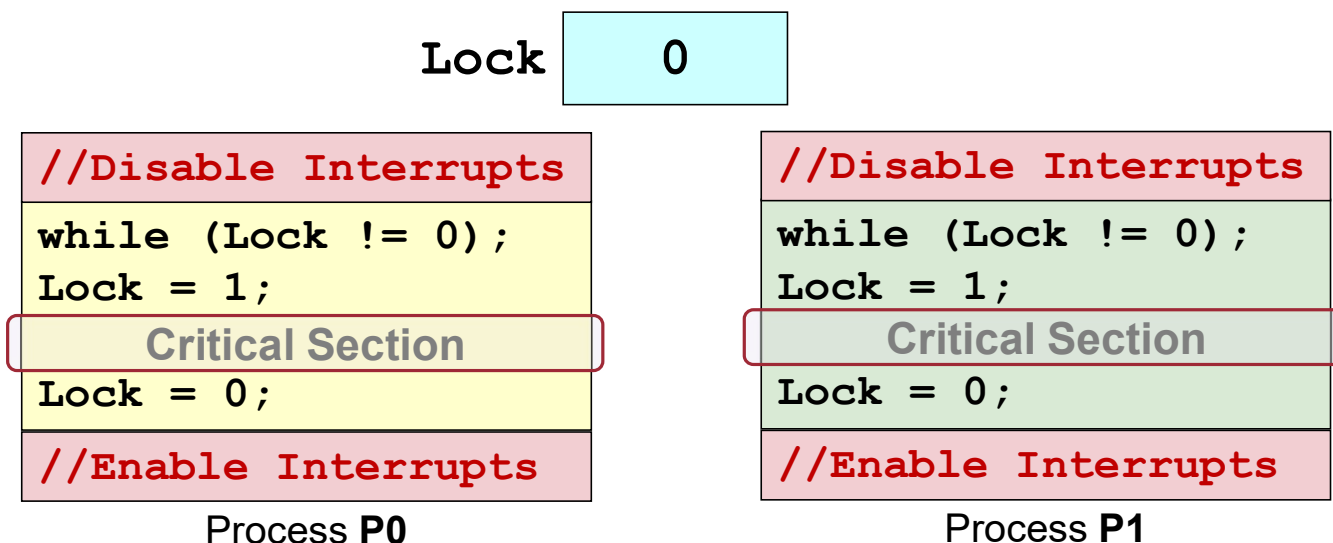
Progress:

Bounded Wait:

Independence:

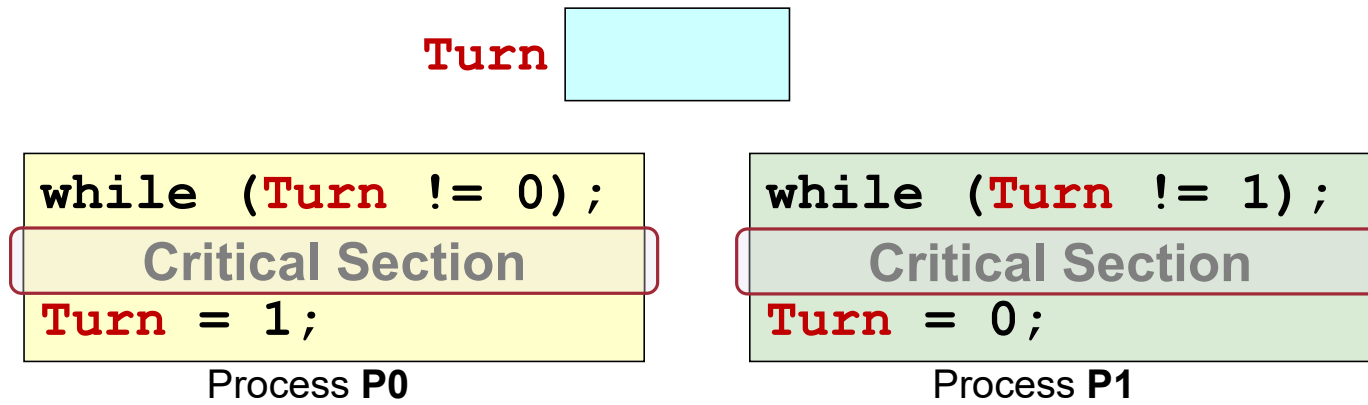
- Makes intuitive sense 😊
 - ❑ But it doesn't work properly ☹️
- It violates the “Mutual Exclusion” requirement!
 - ❑ How?

Using HLL : **Attempt 1 Fixed***



- Solves the problem by preventing context switch
- However:
 - ❑ Buggy critical section may stall the WHOLE system
 - ❑ Busy waiting
 - ❑ Requires permission to disable/enable interrupts

Using High Level Language: **Attempt 2**



■ **Assumption:**

- ❑ P0 and P1 executes the above in loop
- ❑ Take turn to enter **critical section**

■ **Problems:**

- ❑ Starvation:
 - E.g. If P0 never enters CS, P1 starve
- ❑ Violates the **independence** property!

Mutual Exclusion:

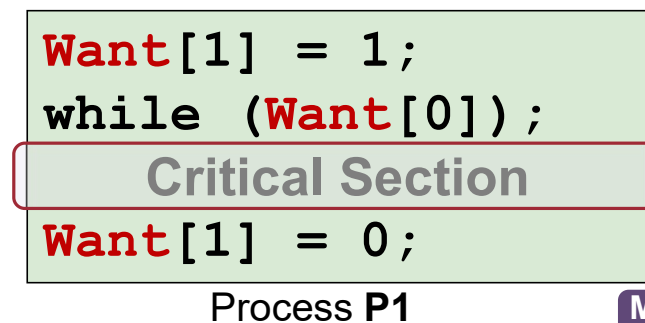
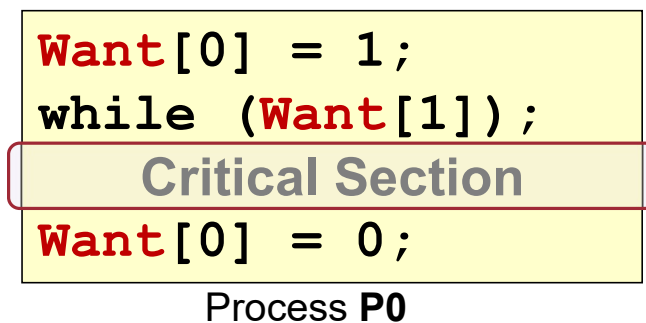
Progress:

Bounded Wait:

Independence:

Using High Level Language: **Attempt 3**

Want[0]
Want[1]



Solves the independence problem

- If P0 or P1 is not around, another process can still enter CS

Problem:

- Deadlock! Try identify the execution sequence that causes deadlock

Mutual Exclusion:

Progress:

Bounded Wait:

Independence:

Using High Level Language: **Attempt 4**



| | |
|-------------|--|
| Turn | |
| Want[0] | |
| Want[1] | |

```
Want[0] = 1;  
Turn = 1;  
while (Want[1] &&  
       Turn == 1);
```

Critical Section

```
Want[0] = 0;
```

Process **P0**

```
Want[1] = 1;  
Turn = 0;  
while (Want[0] &&  
       Turn == 0);
```

Critical Section

```
Want[1] = 0;
```

Process **P1**

Peterson's Algorithm: **Disadvantages**

■ **Busy Waiting:**

- ❑ The waiting process repeatedly test the while-loop condition instead of going into blocked state

■ **Too Low-level:**

- ❑ Higher-level programming construct is desirable
 - Simpler and less error prone

■ **Not general:**


- ❑ General synchronization mechanism is desirable
 - Not just mutual exclusion

Don't worry! The processor has all the answers!


ASSEMBLY LEVEL IMPLEMENTATION

Inspiration: **HLL Attempt 1**

Lock 0



Process **P0**



Process **P1**

Test and Set: An Atomic Instruction

- Machine instruction to aid synchronization
 - Commonly found in modern processors

```
TestAndSet Register, MemoryLocation
```

- **Behavior:**

1. Load the current content at **MemoryLocation** into **Register**
 2. Stores a **1** into **MemoryLocation**
- Important: The above is performed as a **single atomic machine operation**
 - Even in multi-core systems!

Using Test and Set

```
void EnterCS( int* Lock )
{
    while( TestAndSet( Lock ) == 1 );
}
```

\$reg



```
EnterCS ()
    X = X + 1000
ExitCS ()
```

Process P0

\$reg



```
EnterCS ()
    X = X + 1000
ExitCS ()
```

Process P1

Lock



Assume a wrapper function in C that uses test and set:

- takes a memory address M:
- returns the current content at M
- sets content of M to 1

Mutual
Exclusion:

Progress:

Bounded Wait:

Independence:

```
void ExitCS( int* Lock )
{
    *Lock = 0;
}
```

Observations and Comments

- The implementation **works!**
 - However, it employs **busy waiting**
 - Keeps checking the condition until it's safe to enter
 - ➔ Wasteful use of processing power
 - Does not guarantee bounded-wait out of the box:
 - **Unless the scheduling is fair**
 - ➔ But there are algorithms based on test-and-set that address this
- Variants of this instruction exists on most processors:
 - **Compare and Exchange**
 - **Atomic Swap**
 - **Load Link / Store Conditional**

Let's go higher.....

HIGH-LEVEL ABSTRACTION

High Level Synchronization Mechanism

■ Semaphore:

- ❑ A generalized synchronization mechanism
- ❑ **Only functional behavior is specified**
 - ➔ can have different implementations
- ❑ Provides means to:
 - **block** a number of processes
 - ❑ Known as **sleeping processes**
 - **unblock/wake up** one or more sleeping process

■ History:

- ❑ Proposed by **Edgar W. Dijkstra** in 1965

Semaphore: **Wait**() and **Signal**()

- A semaphore **S** contains an integer value
 - Can be initialized to any non-negative values initially
- Two atomic operations:

□ **Wait**(**S**)

- If **s** \leq 0, **blocks** (go to sleep)
- **Decrement** **s**
- Also known as **p()** or **Down()**

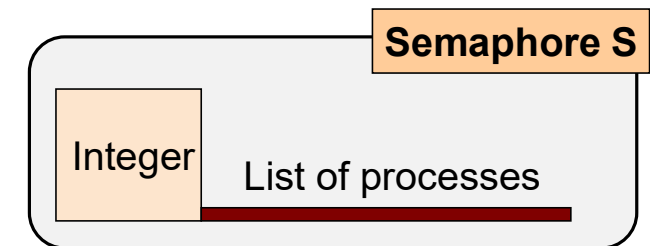
□ **Signal**(**S**)

- **Increments** **s**
- Wakes up one sleeping process if any
- This operation **never** blocks
- Also known as **v()** or **up()**

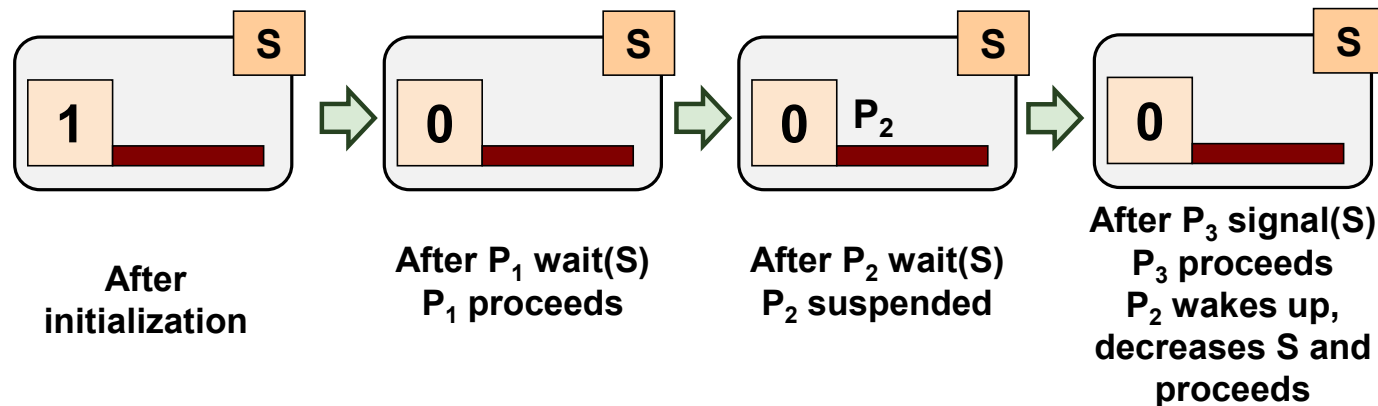
Reminder: The above specifies the **behavior**, not the implementations

Semaphore: **Visualization**


- To aid understanding, you can visualize semaphore as:
 - ❑ A protected integer
 - ❑ A list to keep track of waiting processes



- Example:



Safe Distancing Problem: Only 3 people allowed



- Initially semaphore = 3

These people represent **waiting threads**.
They aren't running on any CPU core.

The bouncer represents a **semaphore**.
He won't allow threads to proceed
until instructed to do so.



- Now semaphore = 0
as 3 people are already in
- New people are waiting
in the queue

Semaphores: **Properties**

- Given:

- $S_{\text{Initial}} \geq 0$

- Then, the following **invariant** must be true:

$$S_{\text{current}} = S_{\text{Initial}} + \# \text{signal}(S) - \# \text{wait}(S)$$

#signal(S) :

- number of signals() operations executed

#wait(S) :

- number of wait() operations **completed**

General and Binary Semaphores

- **General semaphore S:**

- $S \geq 0$ ($S = 0, 1, 2, 3, \dots$)
- also called **counting semaphores**

- **Binary semaphore S:**


- $S = 0$ or 1

- General semaphore is provided for convenience

- **Binary semaphore is sufficient**
- i.e., general semaphore can be mimicked by binary semaphores

Semaphore Example: **Critical Section**

- Binary semaphore **S** = 1



- In this case, **S can only be 0 or 1**
 - From the semaphore invariant
- This semaphore is commonly known as **mutex** (**mut**ual **ex**clusion)

Mutex: Correct CS - Informal Proof

■ Mutual Exclusion:

- N_{CS} = Number of process in critical section
= Process that completed `wait()` but not `signal()`

$$N_{CS} = \#Wait(S) - \#Signal(S)$$

- $S_{Initial} = 1$
- $S_{current} = 1 + \#Signal(S) - \#Wait(S)$

$$S_{current} + N_{CS} = 1$$

- Since $S_{current} \geq 0 \Rightarrow N_{CS} \leq 1$

```
Wait( s );
```

```
Critical Section
```

```
Signal( s );
```

Mutex: Correct CS - Informal Proof (cont)


■ Deadlock:

- Deadlock means **all** processes stuck at `wait(S)`

→ $S_{\text{current}} = 0$ and $N_{\text{CS}} = 0$

$$S_{\text{current}} + N_{\text{CS}} = 1$$

- → ← (contradiction)

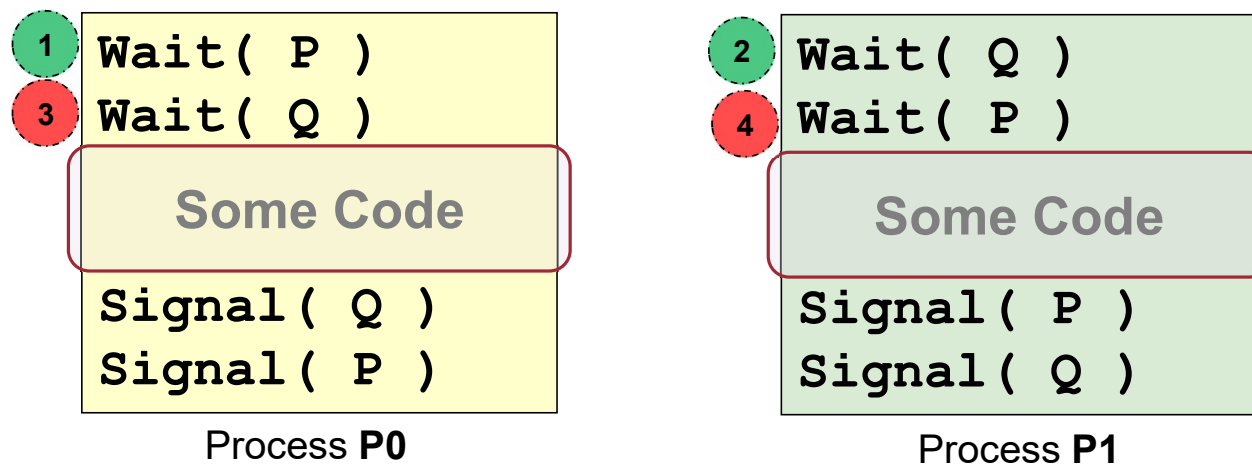


■ Starvation:

- Suppose **P1** is blocked at `wait(S)`
- **P2** is in CS, exits CS with `signal(S)`
 - a. If no other process sleeping, **P1** wakes up
 - b. If there are other process, **P1** eventually wakes up (under fair scheduling)

Incorrect Use of Semaphore: **Deadlock**

- Deadlock is still possible with **incorrect use of semaphore**
- Example: Semaphore **P = 1**, **Q = 1** initially



Semaphore as General Synchronization Tool

- Execute B in P_1 **only after** A executed in P_0
- Use semaphore s initialized to 0

■ Code:


| $P0$ | $P1$ |
|-------------|-----------|
| . | . |
| . | . |
| . | . |
| A | $wait(s)$ |
| $signal(s)$ | B |
| . | . |

Summary: Most Common Uses Of Semaphores

Critical Section

```
BinarySemaphore mutex = 1;


void use_restroom() {
    wait(mutex);
    critical_section();
    signal(mutex);
}
```



Safe-distancing problem

```
Semaphore sem = N;

void client() {
    wait(sem);
    safe_eating();
    signal(sem);
}
```



General synchronization

```
Semaphore sem = 0;
```

```
Process P1:
    produce(X);
    signal(sem);
```

```
process P2:
    wait(sem);
    consume(X);
```

Other High-Level Abstractions

- Semaphore is **very powerful**:
 - ❑ There are no known unsolvable synchronization problem with semaphore (so far 😊)
 - ❑ Other high-level abstractions essentially provide extended features
 - Usually features that are troublesome to express using semaphore alone
- Common alternative: **Conditional Variable**
 - ❑ Allow a task to wait for certain event to happens
 - ❑ Has the ability to ***broadcast***, i.e., wakes up all waiting tasks
 - ❑ related to **monitors**

POSIX Semaphore

- Popular implementation of semaphores in Unix
- Header File:
 - ❑ `#include <semaphore.h>`
- Compilation Flag:
 - ❑ `gcc something.c -lrt`
 - ❑ Stand for "real time library"
- Basic Usage:
 - ❑ Initialize a semaphore
 - ❑ Perform `wait()` or `signal()` on semaphore

pthread Mutex and Conditional Variables

- Synchronization mechanisms for pthreads
- Mutex (`pthread_mutex`):
 - ❑ Binary semaphore (i.e. equivalent `Semaphore(1)`).
 - ❑ Lock: `pthread_mutex_lock()`
 - ❑ Unlock: `pthread_mutex_unlock()`
- Conditional Variables(`pthread_cond`):
 - ❑ Wait: `pthread_cond_wait()`
 - ❑ Signal: `pthread_cond_signal()`
 - ❑ Broadcast: `pthread_cond_broadcast()`

Programming language Support

- Programming languages with thread support must have some synchronization mechanisms
- Examples:
 - ❑ **Java**: all object has built-in lock (mutex), **synchronized** method access, monitors, etc.
 - ❑ **Python**: supports mutex, semaphores, conditional variables, etc.
 - ❑ **C++**: Added built-in thread in C++11; Support mutexes, conditional variable

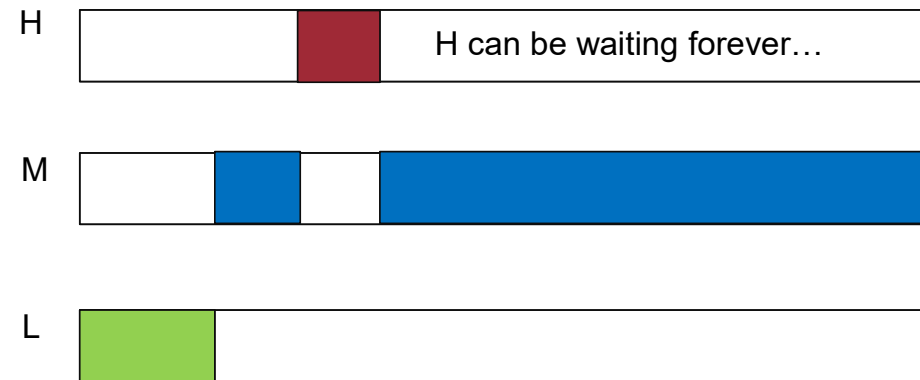
Recall **Priority Scheduling**: Priority Inversion

- Consider the scenario:

- Priority: {H = 1, M=3, L= 5} (1 is highest)
- **L** starts and **locks** a resource (e.g., a file)

- **M** arrives and preempts **L** (higher priority)
 - **L** is unable to unlock the file

- **H** arrives and need the same resource as **L**
 - but the resource is locked! And L can't run because of M
- **M** continues executes even if Task **H** has higher priority
- **M** can starve **H** for as long as it wants!




- Known as **Priority Inversion**:

- Lower priority task effectively preempts higher priority task

Priority Inversion: **Solution**

Priority inheritance

- Temporarily increase priority of **L** to **H**
 - Until it unlocks the lock
- Low-priority job *inherits* the priority of the higher priority process
 - Happens when the high-priority process requests a lock held by low-priority process
 - Priority restored upon successful unlock()



Summary

Synchronization:

- Problem: Race conditions
- Solution: Critical Section
- Criteria for a good solution:
 - ❑ Mutual Exclusion,
 - ❑ progress,
 - ❑ bounded waiting time,
 - ❑ independence
- Important High-Level Construct: Semaphore
 - ❑ Examples in Unix

References

- Modern Operating System (4th Edition)
 - ▣ Chapter 2.5
- Operating System Concepts (9th Edition)
 - ▣ Chapter 5
- Three Easy Pieces
 - ▣ Chapters 25, 26,, 34!
- Edgar W. Dijkstra, “Note No.123: Cooperating Sequential Processes”
 - ▣ <http://www.cs.utexas.edu/users/EWD/ewd01xx/EWD123.PDF>